AZURE, DEVOPS, KUBERNETES

Interacting with OpenShift Local in an Azure VM

29.11.2023 by **DevOps**

In my previous **blog** I've written about the no or low cost solutions to setup an OpenShift cluster to play with. I've detailed the installation of OpenShift Local on Azure. OpenShift Local is like a mini cluster that is installed on just one virtual machine (node). I've also shared the tips on how to properly stop the cluster before shutting done the Azure VM. In this blog I'll share with you how you can connect to your cluster from your laptop. I'll also share more tips & tricks I've gathered along the way that may be good to know.

Starting the Azure VM and the

able to run commands to interact with our cluster. However this access is only command line and it would be nice to be able to connect to the OpenShift Web console and benefit fully from what we can do with OpenShift.

At this stage you may have stopped your Azure VM and so let's start it again and ssh to it when it is running. As I said in my previous blog, you have to check the content of the file /etc/resolv.conf as it may change. I'll explain more on that below but at this stage the content may be the same as the one you've noted when you've installed the cluster. Otherwise, just put back the original content. This is important because if the content has been changed, you'll get the following DNS error when you try to start the cluster:

```
INFO Check DNS query from host...
WARN foo.apps-crc.testing resolved to [x.x.x.x] but 1
```

The DNS resolution point to the IP Address (x.x.x.x) of your interface eth0 instead of the fixed IP Address of the OpenShift Local instance (192.168.130.11) and the cluster will not start.

When this DNS topic is cleared we can then start our OpenShift cluster:

```
$ crc start
```

I didn't mentioned it previously but **crc** is the embedded tool used to manage our OpenShift Local cluster. As during our installation, starting the cluster takes several minutes (around 10 minutes) and at the end you will get your credentials to connect to it. Let's now see how to configure our Azure VM to be able to connect directly from our laptop.

Connect to OpenShift Local from your laptop

We need to configure three components in our Azure VM: Firewalld, HA Proxy and NetworkManager. Let's start wit firewalld:

```
[enb@enb-OpenShift ~]$ sudo firewall-cmd --add-port=
[enb@enb-OpenShift ~]$ sudo firewall-cmd --add-port=
[enb@enb-OpenShift ~]$ sudo firewall-cmd --add-port=
[enb@enb-OpenShift ~]$ sudo systemctl restart firewa
[enb@enb-OpenShift ~]$ sudo firewall-cmd --list-all
public (active)
  target: default
  icmp-block-inversion: no
  interfaces: eth0
  sources:
  services: cockpit dhcpv6-client ssh
  ports: 80/tcp 6443/tcp 443/tcp
  protocols:
  forward: no
  masquerade: no
  forward-ports:
  source-ports:
  icmp-blocks:
  rich rules:

[enb@enb-OpenShift ~]$ sudo semanage port -a -t http
```

So you just open the required port in the operating system firewall to connect to the cluster. In my previous blog, you may remember that we also opened those ports in the NSG of the public interface of our Azure VM. You now have access to those ports to connect directly to your cluster.

The next step is to install and configure HA Proxy which is used to receive the request from your laptop on those ports and direct them to the cluster:

```
[enb@enb-OpenShift ~]$ sudo dnf -y install haproxy p
[enb@enb-OpenShift ~]$ sudo cp /etc/haproxy/haproxy.

[enb@enb-OpenShift ~]$ tee haproxy.cfg &>/dev/null <
global
        debug

defaults
        log global
        mode    http
        timeout connect 5000
        timeout client 5000
        timeout server 5000

frontend apps
    bind SERVER_IP:80
    bind SERVER_IP:443
    option tcplog
    mode tcp
    default_backend apps

backend apps
    mode tcp
    balance roundrobin
    option ssl-hello-chk
    server api.crc.testing CRC_IP:443 check

frontend api
    bind SERVER_IP:6443
    option tcplog
    mode tcp
    default_backend api

backend api
    mode tcp
    balance roundrobin
    option ssl-hello-chk
    server api.crc.testing CRC_IP:6443 check
EOF

[enb@enb-OpenShift ~]$ export SERVER_IP=$(hostname -
[enb@enb-OpenShift ~]$ export CRC_IP=$(crc ip)
[enb@enb-OpenShift ~]$ sed -i "s/SERVER_IP/$SERVER_I
[enb@enb-OpenShift ~]$ sed -i "s/CRC_IP/$CRC_IP/g" h
[enb@enb-OpenShift ~]$ sudo cp haproxy.cfg /etc/hapr
[enb@enb-OpenShift ~]$ sudo systemctl start haproxy
[enb@enb-OpenShift ~]$ sudo systemctl enable haproxy
```

After downloading the haproxy package, you create its configuration file and fill it with the appropriate IP Addresses of your VM as well as the cluster. At this stage HA Proxy is up and running!

The last component to configure is the NetworkManager of the Azure VM:

```
[enb@enb-OpenShift ~]$ sudo tee /etc/NetworkManager/
[main]
dns=dnsmasq
EOF

[enb@enb-OpenShift ~]$ tee external-crc.conf &>/dev/
address=/apps-crc.testing/SERVER_IP
address=/api.crc.testing/SERVER_IP
EOF

[enb@enb-OpenShift ~]$ export SERVER_IP=$(hostname -
[enb@enb-OpenShift ~]$ sed -i "s/SERVER_IP/$SERVER_I
[enb@enb-OpenShift ~]$ sudo cp external-crc.conf /et
[enb@enb-OpenShift ~]$ sudo systemctl reload Network
```

We've configured NetworkManager to use dnsmasq and this is where the file /etc/resolv.conf is changed when you reload NetworkManager (or restart the VM). This configuration is OK as long as the cluster is running, however when you stop and restart it, you need to recover the initial configuration of /etc/resolv.conf as mentioned above.

Setting on your laptop

On your laptop you have to add the following in the file /etc/hosts (or \WINDOWS\system32\drivers\etc\hosts if you are using Windows) where x.x.x.x is the public IP Address of your Azure VM:

```
x.x.x.x api.crc.testing canary-openshift-ingress-cana
```

You have now an access to the OpenShift Web console when you enter **https://console-openshift-console.apps-crc.testing**

in your web browser. Choose insecure connection when asked and you will see the login screen:

You can also use now the command line interface directly from your laptop:

```
% oc login -u kubeadmin -p <kubeadmin_password> https
```

This is convenient if you want to apply some yaml file you have on your laptop directly to the cluster without the need to copy them to the Azure VM first.

Congratulations! You are now fully ready to test your OpenShift Local cluster!

Last words to keep your environment healthy: When you need to shut your Azure VM down, proceed as recommended in my previous blog first:

```
$ oc logout
$ crc stop
```

Then stop the VM. When you start it up again, proceed as follows after restoring first the /etc/resolv.conf file:

```
$ crc start
$ crc status
```

When crc start is completed, the cluster may not be immediately available so check its status with crc status and wait to see **OpenShift: Running**.

That's all for today! I wish you much fun with OpenShift!

Post Views: 2,141

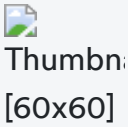

by
DevOps

No comments yet.

Leave a Reply:

Email

name@example.com

Username

Comment

hCaptcha

Submit Cancel

BE SHARING

Related blog articles

ANSIBLE, DEVOPS

Parallel execution of Ansible roles

10.06.2025 by **Martin Bracher**

AZURE, CLOUD

Azure Bootcamp Switzerland 2025

10.06.2025 by **Nicolas Jardot**

AZURE

Azure Bootcamp Switzerland 2025, an Azure day in Bern

06.06.2025 by **Adrien Devaux**

AZURE, BUSINESS INTELLIGENCE, CLOUD, DATABASE ADMINISTRATION & MONITORING, DATABASE MANAGEMENT

SQLDay 2025 – Wrocław – Sessions

19.05.2025 by **Amine Haloui**

ABOUT

dbi services is a company specialized in IT consulting and services. We are experts in innovative and efficient data infrastructures and platforms. Tailor-made solutions is what we offer to our customers thanks to our consultants, whose skills and knowledge are constantly evolving thanks to continuous training.

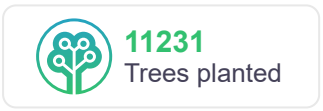

WE ARE CLOSE TO YOU

Delémont >
Nyon >
Bâle >
Berne >
Zurich >